SYSTEMS ARCHITECTURE & GREEN MULTI-AGENT PROTOCOLS · CANONICAL ARCHIVAL RELEASE V5.4.5

TOAI-elegant-lang: A Tri-Script Orthogonal RISC Wire Protocol, Adversarial Corpus Audit, and Empirical Comprehension Harness for Multi-Agent Communication

Koji Toriyama (PhenoX) · Principal Systems Architect, PhenoX AI Alliance, Kawasaki, Japan

Repository & Reproducibility Artifact: https://github.com/PhenoX-AI-Alliance/TOAI-elegant-lang | **PyPI:** `pip install toai-elegant-lang==5.4.5`

ABSTRACT

Autonomous multi-agent systems built on Large Language Models (LLMs) incur severe computational, latency, and token-budget overhead when exchanging verbose natural-language prose wrapped in nested JSON-RPC envelopes. Compressing agent-to-agent (A2A) messages exposes a fundamental systems trilemma among *wire token efficiency*, *syntactic, numeric, and causal execution fidelity*, and *human auditability*. This paper presents **TOAI-elegant-lang (v5.4.5)**, a zero-dependency, deterministic compiler and wire protocol inspired by the Rosetta Stone that aligns multi-agent state transitions across three synchronized representations: (1) **Demotic Wire Script**, a compact causal wire syntax (`@CTX:... :: §S01 ... §X64`); (2) **Sacred Expanded Script**, a deterministic macro- and primitive-expanded canonical form (`§GAIA_WARMTH_HARMONY`); and (3) the **Rosetta Audit Layer**, which attaches natural-language definitions and domain invariants for human inspection. Under `o200k_base` and `cl100k_base` Byte-Pair Encoding (BPE) tokenizers, each canonical §-prefixed identifier strictly occupies 3 BPE tokens ('§' + `Prefix` + `Digits`), while the 256-primitive + 256-macro codebook fits inside a $\sim$1,000-token system prompt header. Evaluated across standard A2A/MCP JSON serializers (`benchmarks/fair_bench.py`), Demotic Wire averages **73.2 BPE tokens/message**, achieving a **79.1% token reduction vs. Verbose A2A JSON** (349.8 tokens) and **55.7% vs. Structured Tool JSON** (165.2 tokens) while maintaining **100.0% downstream AST/task execution accuracy** (outperforming `LLMLingua`-style perplexity pruning at $\rho = 0.35$, where heuristic task recovery degrades to 61.6%), remaining fully orthogonal to state-delta encoding, and breaking even within **4–20 messages uncached** or **1–2 messages under 90% prompt caching**. Following multi-AI adversarial auditing, we introduce a non-circular **Independent Corpus Alignment Auditor** (`src/audit_corpus_alignment.py`) validated against a human gold-standard sample ($n = 60$, $p_o = 93.33\%$, **Cohen's** $\kappa = 0.8611$). To eliminate circularity, receiver comprehension is evaluated on a **strictly disjoint** $N = 112$ **Holdout Benchmark** (0.0% wire/scenario overlap) across **7 frontier LLMs** from Google, OpenAI, Anthropic, and DeepSeek: bare wire collapses to chance level (**25.1% mean**, matching 25.0% random chance), while expanded identifiers restore semantic comprehension (**77.3% mean**). Crucially, we identify the **Flagship Reasoning Paradox**: deep reasoning models (Claude Sonnet 5.5) overthink abbreviated syntax and inflate internal thinking tokens, scoring lower than fast models that rely on direct semantic matching. This establishes that compressed symbolic notation is practically bounded to **lightweight worker RPC queues**, not flagship reasoning deliberation. Finally, origin attestation is formalized with a game-based EUF-CMA reduction bounding k-turn internal-counter forgery by $10^{-4k} + \epsilon_{\text{PRF}} + \epsilon_{\text{coll}}$.

1. Introduction and Multi-Agent Communication Trilemma

As Large Language Model (LLM) applications evolve from single-turn assistants into multi-agent workflows—e.g., Anthropic Model Context Protocol (MCP) [1], Google Agent-to-Agent (A2A) JSON-RPC 2.0 orchestrators [2], and AutoGen [5]—inter-agent communication overhead has become the dominant bottleneck in latency, context saturation, and KV-cache footprint. In conventional topologies, agents coordinate state transitions by generating conversational prose inside nested JSON schemas. While human-readable, schema keys and punctuation delimiters waste token budgets and scale attention FLOPs quadratically. Compressing inter-agent messages exposes a strict **Multi-Agent Communication Trilemma**:

1. **Wire Token Efficiency:** Minimizing BPE subword tokens per transition and amortizing shared schemas across sessions.
2. **Syntactic, Numeric, and Causal Execution Fidelity:** Preserving structured parameters (e.g., leading-zero strings `"007"`, embedded commas `"a, b"`, floats `150.50`), transition probabilities $p \in [0, 1]$, costs $c \in [0, 0.9999]$, and fault edges (`!->`) for downstream AST execution.
3. **Human Auditability and Cross-Model Portability:** Allowing humans and heterogeneous LLM families (GPT, Claude, Gemini, Llama) to inspect and exchange packets over standard text APIs without opaque binary blobs or model-specific latent adapters.

Existing approaches sacrifice at least one vertex of this trilemma (Table 1). Perplexity-based prompt compression (LLMLingua [4]) drops tokens dynamically but breaks JSON syntax delimiters and distorts numeric thresholds at high compression ratios ($\rho \leq 0.35$, Section 5). Emergent discrete MARL codes [6] or continuous latent vectors [7] sacrifice human auditability and fail across commercial black-box APIs.

Recent structured multi-agent protocols address related aspects but occupy distinct architectural niches:

- **Graph-Grounded Protocols (G2CP):** Replaces free-form dialogue with explicit knowledge graph operations, achieving $\approx 73\%$ token savings. While G2CP excels in graph traversal tasks, it requires external graph database backends; TOAI-elegant-lang provides general-purpose causal AST exchange without infrastructure dependencies.
- **KV-Cache Transmission (Q-KVComm / KVComm):** Transmits quantized Key-Value cache activations directly across model layers (5–6× compression). This requires white-box model weights and homogeneous architectures; TOAI operates strictly at the application-layer text wire across heterogeneous black-box commercial APIs.
- **Dynamic Topology Pruning (AgentDropout):** Eliminates redundant agent nodes and communication edges. This optimizes workflow topology and is fully orthogonal to wire serialization: AgentDropout reduces *who* talks to *whom*, while TOAI minimizes the wire payload *when* they talk.
- **State Delta Encoding (SDE):** Transmits attribute diffs rather than full state objects. Because TOAI's typed parameter brackets (`[key:val]`) are schema-agnostic, state deltas can be nested directly inside Demotic payloads (e.g., `§S01[delta:{+limit:10}]`), providing synergistic, orthogonal token reduction without serialization overhead.

To address this trilemma, we present **TOAI-elegant-lang (v5.4.5)**, a zero-dependency Python compiler and wire protocol built on an orthogonal Reduced Instruction Set Computer (RISC) architecture and Tri-Script Rosetta design. Crucially, release `v5.4.5` incorporates multi-AI adversarial audits (ChatGPT and Claude) and rigorous peer evaluations, replacing provisional estimates with a strictly disjoint $N = 112$ holdout benchmark, complete downstream execution accuracy comparisons, and a formal game-based cryptographic reduction.

Table 1: Architectural Comparison of Multi-Agent Communication Paradigms (Verified Standards & Repository Baselines)

Communication Paradigm	Wire Encoding	Causal AST	Deterministic	Portability	Auditability & Operational Scope
Verbose A2A / MCP JSON [1, 2]	Pretty JSON + Prose	Yes (Verbose)	100.0%	100% (Text API)	Full readability; highest token overhead (349.8 tok/msg).
Structured Tool JSON	Standard Schema	Yes (JSON)	100.0%	100% (Text API)	Lossless schema; repeated key boilerplate (165.2 tok/msg).
G2CP (Graph-Grounded)	Subgraph Ops	Graph Only	100.0%	100% (Text API)	Task-verifiable; requires dedicated Knowledge Graph backend.
Q-KVComm / KVComm	Quantized KV-Cache	Internal	Lossy (Cosine)	White-Box Only	High compression (5×); cannot span heterogeneous black-box APIs.
AgentDropout	Graph Pruning	Topology	N/A (Topology)	Orthogonal	Prunes inactive communication links; orthogonal to serialization.
State Delta Encoding (SDE)	Attribute Diffs	Diff Objects	100.0%	100% (Text API)	Efficient for stateful updates; orthogonal to TOAI wire brackets.
TOAI-elegant-lang (Ours)	**Demotic RISC Wire**	**Yes (LL(1) AST)**	**100.0% Lossless**	**100% (Text API)**	**Tri-Script Rosetta; 73.2 tok/msg; zero runtime dependencies.**

Orthogonal Composition with State-Delta Encoding (SDE): Because State-Delta Encoding (SDE) eliminates redundant

past-turn state keys from conversation histories while TOAI compresses the wire syntax of active state transitions, the two techniques compose multiplicatively: embedding Demotic Wire within an SDE delta envelope reduces message payload from an uncompressed 349.8 BPE tokens to **29.3 BPE tokens** (an **83.2% compound token reduction**), verifying full practical orthogonality without AST deserialization loss.

2. System Architecture, LL(1) Grammar, and Quote-Aware Stateful Codec

2.1 Tri-Script Architecture and 8-Bit Orthogonal RISC Lexicon Hierarchy

TOAI-elegant-lang decouples machine transmission from human inspection via three synchronized layers (Table 2):

1. **Demotic Wire Script (Compact Machine Layer):** Transmitted between agents. Combines context header (`@CTX:DOMAIN`), 1-byte logical core opcodes (`§S01`–`§H32` across 8 categories: System, Data, Governance, Fault, Execution, Audit, Gaia, Logic; exactly 256 primitives in `core_lexicon.json`), 256 active higher-order macros (`§X01`–`§X969` in `macro_lexicon.json`; under both `o200k_base` and `cl100k_base`, each canonical identifier strictly tokenizes into 3 BPE tokens, e.g., `§X64` → `["§", "X", "64"]` and `§S01` → `["§", "S", "01"]`), typed payloads (`[k:v]`), and causal operators (`-[p:..., c:...]->`, `=>`, `!->`).
2. **Sacred Expanded Script (Canonical Symbolic Layer):** Deterministic expansion via `codec.decode(wire)` expanding opcodes to canonical uppercase snake-case identifiers matching strict `^§[A-Z][A-Z0-9_]*$` (e.g., `§GAIA_WARMTH_HARMONY`).
3. **Rosetta Audit Layer (Human Governance Layer):** Maps every symbol to concise natural-language definitions for compliance auditing (`rosetta_layer.json`) and Condition C glossary injection.

Exact BPE Token Accounting: Under `o200k_base` and `cl100k_base`, the section sign `'§'` (`U+00A7`) is tokenized separately from following ASCII characters. In all reported benchmarks, every canonical §-prefixed opcode (`§S01`, `§G48`, `§X64`) is strictly counted as **3 BPE tokens** ('§' + `Prefix` + `Digits`), confirming that our 73.2 tok/msg average includes full token penalties.

Table 2: Tri-Script Representation and Lexicon Hierarchy in TOAI-elegant-lang v5.4.5

Architectural Layer	Source / Capacity	Canonical Pattern / Example	Validation & Runtime Role in v5.4.5
1. Demotic Wire Script (Core Primitives)	`core_lexicon.json` (256 Primitives)	`§S01..§L32`, `§G48`, `§C20` (3 BPE tok/symbol)	Immutable 256-primitive base vocabulary; validated at startup and by `test_17` for bijective mapping.
1b. Demotic Wire Script (Active LRU Macros)	`macro_lexicon.json` (256 Active + Archive)	`§X01..§X969` (e.g., `§X64`, `§X354`; 3 BPE tok/symbol)	Depth-1 macro expansion into core primitives; pinned core macros prevent LRU thrashing.
2. Sacred Expanded Script (Canonical AST)	`codec.decode()` Output	`^§[A-Z][A-Z0-9_]*$` (e.g., `§GAIA_WARMTH_HARMONY`)	Zero lowercase or CJK corruption permitted; strictly enforced by unit test `test_17` across all files.
3. Rosetta Audit Layer (Human Definitions)	`rosetta_layer.json` & `definitions`	Natural-language semantic specification per symbol	Used for human governance inspection, `audit_corpus_alignment.py`, and `eval_comprehension.py --cond C`.
Inline Structured Payload (`[k:v]`)	Quote-Aware Scanner (`_parse_payload_block`)	`[pair:"USD/JPY", stop:150.5, code:"007", note:"a, b"]`	Stateful character scanner preserving quoted strings, leading zeros, commas, brackets, and escapes.

Orthogonal Composition with State-Delta Encoding (SDE): Because State-Delta Encoding (SDE) eliminates redundant past-turn state keys from conversation histories while TOAI compresses the wire syntax of active state transitions, the two techniques compose multiplicatively: embedding Demotic Wire within an SDE delta envelope reduces message payload from an uncompressed 349.8 BPE tokens to **29.3 BPE tokens** (an **83.2% compound token reduction**), verifying full practical orthogonality without AST deserialization loss.

2.2 Conflict-Free LL(1) EBNF Grammar and Quote-Aware Stateful Payload Parser

To guarantee deterministic $\mathcal{O}(L)$ parsing without regex backtracking bugs or payload corruption, `v5.4.5` implements a conflict-free LL(1) grammar paired with a stateful character-level quote/escape scanner in `src/codec.py`:

```
packet = [ "@CTX:", identifier, "::" ], expr_node, { ( causal_arrow | "=>" | "!->" | ">" | "?" ), expr_node } ;
expr_node = ( "(", packet, ")" | "{", packet, "}" | "§", ( risc_id | macro_id | sacred_name ) ), [ "[", kv_pair, { ",", kv_pair }, "]" ] ;
kv_pair = identifier, ":", ( '"', { ( utf8_char - ( '"' | '\' ) ) | '\"' | '\\' }, '"' | float_lit | int_lit | bool_null_id ) ;
causal_arrow = "-[", [ "p:", float_val, "," ], ( "c:" | "cost:" ), cost_mantissa, "]->" ; (* cost_mantissa: exact 4 digits *)
```

The parser strictly distinguishes quoted strings (`"007"` → `str`) from unquoted numbers (`007` → `int`), isolates delimiters inside quotes (`note:"a, b"`, `expr:"x]y"`), and enforces explicit OOV policies (`on_oov="error" | "keep" | "nearest"` with an audit log in `EncodedWireResult.substitutions`, verified by tests `test_07`–`test_15`).

3. Formal Threat Model and Game-Based EUF-CMA Proof of Origin Attestation

3.1 Unobservable Internal-Counter Attestation Construction (Zero Timestamp Jitter)

To prevent an active Dolev-Yao adversary $\mathcal{A}$ from injecting forged commands or observing rejection-sampling loops under an adaptive Chosen-Message Attack (EUF-CMA), TOAI-elegant-lang embeds an origin tag directly into the 4-digit causal cost mantissa $c_n \in [0.0000, 0.9999]$. Session key $S_0 \in \{0,1\}^{256}$ is established via authenticated X25519 ECDH (with static host keys) and HKDF-SHA256, guaranteeing cross-session replay immunity; strict timestamp monotonicity $t_n > t_{n-1}$ enforces per-turn freshness and rejects clock skew at the receiver. Let $H_0 = \text{SHA256}(\texttt{"TOAI_TELEPATH_GENESIS_V541"} \,\|\, \text{SHA256}(\texttt{core_lexicon.json}) \,\|\, \text{SHA256}(\texttt{macro_lexicon.json}))$, and $H_{n-1} = \text{SHA256}(H_{n-2} \,\|\, \texttt{0x00} \,\|\, \text{wire}_{n-1})$. For monotonic timestamp $t_n > t_{n-1}$, sender and verifier compute the 16-bit value $Y_n \in \{0, \dots, 59{,}999\}$ for public confidence mantissa $P_n = \lfloor p_n \times 10^4 \rfloor \in \{0, \dots, 9999\}$ via an **internal, non-transmitted rejection counter** $r_n \in \{0, 1, \dots\}$:

$$r_n = \min\Big\{ r \in \mathbb{N}_0 : \big(P_n \oplus \text{trunc}_{16}(\text{HMAC-SHA256}(S_0,\, H_{n-1} \,\|\, \texttt{0x00} \,\|\, \text{dom}(t_n, r)))\big) < 60{,}000 \Big\} \tag{1}$$

where $\text{dom}(t_n, 0) = t_n$ and $\text{dom}(t_n, r) = t_n \,\|\, \text{uint8}(r)$ for $r \geq 1$. **Collision-Free Disambiguation:** Because r_n is defined by the strict minimum operator (min), both sender and verifier deterministically halt at the exact same first counter value $r \geq 0$ satisfying $Y_n^{(r)} < 60{,}000$, completely eliminating any ambiguity or collision when multiple higher counters also satisfy the inequality. **Wire Transmission and Receiver Verification:** The wire transmits the pair (P_n, C_n) where $P_n = \lfloor p_n \times 10^4 \rfloor \in \{0, \dots, 9999\}$ is the causal confidence and $C_n = \lfloor c_n \times 10^4 \rfloor \in \{0, \dots, 9999\}$ embeds the origin attestation. The receiver executes deterministic verification by re-evaluating the internal counter starting from $r = 0$:

$$T_n^{(r)} = \text{trunc}_{16}\big(\text{HMAC-SHA256}(S_0,\, H_{n-1} \,\|\, \texttt{0x00} \,\|\, \text{dom}(t_n, r))\big), \quad Y_n^{(r)} = P_n \oplus T_n^{(r)} \tag{2}$$

The receiver finds the first $r_n \geq 0$ such that $Y_n^{(r_n)} < 60{,}000$ (identical to sender evaluation) and accepts the turn if and only if $Y_n^{(r_n)} \bmod 10{,}000 == C_n$. The verifier does *not* invert Y_n; rather, the verifier computes the forward PRF sequence and tests equivalence. Because bitwise XOR is a bijection on $\{0, \dots, 65{,}535\}$, each candidate $Y_n^{(r)}$ is uniform with single-draw acceptance probability $p_{\text{acc}} = 60{,}000/65{,}536 = 91.55\%$. In this geometric trials setting, the expected internal evaluations per turn is $\mathbb{E}[r_n + 1] = 1/p_{\text{acc}} = 65{,}536/60{,}000 \approx 1.09226 \leq 1.0923$, bounding total PRF oracle queries by $\mathbb{E}[Q_{\text{PRF}}] \leq 1.0923(Q + k)$. For $r_1 = 0$ (Table 3), $Y_1 = 9841 \oplus 9247 = 622 < 60{,}000$, yielding $C_1 = 622 \bmod 10{,}000 = 622 \implies c_1 = 0.0622$.

3.2 Formal Security Reduction ($\text{Game}_0 \to \text{Game}_2$) and Congruence Cardinality Proof

Theorem 1 (Exact EUF-CMA Forgery Bound Subsuming Direct Tag Guessing and Modular Aliases). *Let $\mathcal{A}(\tau, Q, k)$ be an adaptive EUF-CMA adversary making at most Q signing queries to MAC oracle $\mathcal{O}_{MAC}(S_0, \cdot)$ and outputting a k-turn unqueried forgery trajectory $\{(wire^*_{n+j}, t^*_{n+j}, P^*_{n+j}, C^*_{n+j})\}_{j=1}^{k}$ with $t^*_{n+j} > t^*_{n+j-1}$. Under the $(\tau', Q_{PRF}, \epsilon_{PRF})$-PRF security of HMAC-SHA256 (where $\mathbb{E}[Q_{PRF}] \le 1.0923(Q+k)$) and (τ', ϵ_{coll})-collision resistance of SHA-256:*

$$\text{Adv}^{\text{EUF-CMA}}_{\text{TOAI}}(\mathcal{A}) \le \left(\frac{6}{60{,}000}\right)^k + \epsilon_{\text{PRF}} + \epsilon_{\text{coll}} = (10^{-4})^k + \epsilon_{\text{PRF}} + \epsilon_{\text{coll}} = 10^{-4k} + \epsilon_{\text{PRF}} + \epsilon_{\text{coll}} \tag{3}$$

Proof. We define three security games $\text{Game}_0 \to \text{Game}_1 \to \text{Game}_2$, letting W_i denote $\mathcal{A}$'s winning event in Game_i:

- **Game_0 (Real EUF-CMA Experiment):** Challenger samples $S_0 \xleftarrow{\$} \{0,1\}^{256}$, answers Q queries via Eqs. (1)–(2), and evaluates $\mathcal{A}$'s k-turn forgery. By definition, $\Pr[W_0] = \text{Adv}^{\text{EUF-CMA}}_{\text{TOAI}}(\mathcal{A})$.
- **Game_1 (Transcript Collision Abort):** Challenger aborts if two distinct wire histories ever produce matching chain digest $H = H'$. Any abort yields a SHA-256 collision within $\tau' \approx \tau + \mathcal{O}(Q+k)$, hence $|\Pr[W_0] - \Pr[W_1]| \le \epsilon_{\text{coll}}$.
- **Game_2 (PRF Replacement):** We replace HMAC-SHA256$(S_0, \cdot)$ with a lazy uniform random function $\mathcal{R} : \{0,1\}^* \to \{0,1\}^{256}$. By PRF reduction over at most Q_{PRF} evaluations, $|\Pr[W_1] - \Pr[W_2]| \le \epsilon_{\text{PRF}}$. **Lemma 1 (Deterministic Input Freshness via Monotonic Verification):** Let $\mathcal{Q} = \{(H_i, t_i)\}_{i=1}^{Q}$ be the set of queries submitted by $\mathcal{A}$ to $\mathcal{O}_{\text{MAC}}$, and let (H^*_{n+j}, t^*_{n+j}) be any forged turn submitted to the verifier. Because the verifier maintains monotonic state $t^*_{n+j} > t^*_{n+j-1} \ge \max_{(H,t)\in\mathcal{Q}} t$ (where in distributed settings with clock skew, t_n is deterministically instantiated as a monotonic logical sequence counter $s_n = s_{n-1}+1$), the counter t^*_{n+j} strictly exceeds all previous query counters. Any attempt by $\mathcal{A}$ to replay a previously queried timestamp $t^* \le t^*_{n+j-1}$ is rejected deterministically by the verifier with probability 1. Furthermore, under the collision resistance of SHA-256 (Game_1), distinct transcripts cannot yield $H^* = H$. Thus, for all $r \ge 0$, the PRF input $X^{(r)}_{n+j} = H^*_{n+j-1} \,\|\, \texttt{0x00} \,\|\, \text{dom}(t^*_{n+j}, r)$ is strictly fresh and was never evaluated by $\mathcal{R}$ prior to the forgery attempt.
- **Residue Class Analysis in Game_2:** For each forged turn $j \in \{1, \dots, k\}$: (1) If $C^*_{n+j} \notin \{0, \dots, 9999\}$, the LL(1) parser rejects deterministically ($\Pr = 0$). (2) If $C^*_{n+j} \in \{0, \dots, 9999\}$, the adversary chooses P^*_{n+j} and C^*_{n+j}. Because $T^{(r)}_{n+j} = \text{trunc}_{16}(\mathcal{R}(X^{(r)}_{n+j}))$ is drawn uniformly at random from $\{0, \dots, 65{,}535\}$ independently of (P^*_{n+j}, C^*_{n+j}) and $T \mapsto P^*_{n+j} \oplus T$ is a bijection, each candidate $Y^{(r)}_{n+j} = P^*_{n+j} \oplus T^{(r)}_{n+j}$ is independently and uniformly distributed. Upon conditioning on acceptance ($Y^{(r)}_{n+j} < 60{,}000$), the resulting value Y_{n+j} is strictly uniform over $\{0, \dots, 59{,}999\}$. The winning condition is $Y_{n+j} \equiv C^*_{n+j} \pmod{10{,}000}$. For any $C^*_{n+j} \in \{0, \dots, 9999\}$ and $m \in \{0, \dots, 5\}$, we have $0 \le C^*_{n+j} + 10{,}000m \le 9{,}999 + 50{,}000 = 59{,}999$. Conversely, for $m \ge 6$, $C^*_{n+j} + 10{,}000m \ge 60{,}000$. Thus, every residue class modulo 10,000 has exactly $60{,}000/10{,}000 = 6$ representatives in $\{0, \dots, 59{,}999\}$. The winning set $\mathcal{W}(C^*_{n+j}) = \{C^*_{n+j} + 10{,}000m : m \in \{0, \dots, 5\}\}$ satisfies $|\mathcal{W}(C^*_{n+j})| = 6$ identically for all choices of C^*_{n+j}:

$$\Pr\left[Y_{n+j} \in \mathcal{W}(C^*_{n+j})\right] = \frac{|\mathcal{W}(C^*_{n+j})|}{60{,}000} = \frac{6}{60{,}000} = 10^{-4} \tag{4}$$

This bound strictly accounts for both Direct Tag Guessing ($m = 0$, probability 1/60,000) and Modular Aliases ($m \in \{1, \dots, 5\}$, probability 5/60,000). Because each turn uses a strictly fresh chain hash H^*_{n+j-1} and monotonic timestamp t^*_{n+j}, the draws are mutually independent across turns, yielding $\Pr[W_2] = (6/60{,}000)^k = 10^{-4k}$. **Side-Channel Scope and Limitations:** We emphasize that origin attestation is strictly evaluated against remote network adversaries attempting existential forgery over the communication wire. Physical, cache-timing, or co-located kernel side-channel attacks against the rejection-sampling loop are explicitly out of scope; deployment in adversarial co-tenancy environments would require constant-time modular arithmetic or blinding countermeasures. □

Table 3: Reproducible Cryptographic Test Vector for Equation (2) (Verifiable via Python `hashlib`/`hmac`, $r_1 = 0$)

Parameter / Verification Step	Canonical Hexadecimal / Decimal Value (Bitwise XOR Precedes mod 10,000)
Secret Seed S_0 (256-bit hex)	`7f83b1657ff1fc53b92dc18148a1d65dfc2d4b1fa3d677284addd200126d9069`
Session Anchor H_0 (SHA256)	`f3c515b25afe4f7ed4ca735b85fbfe50b083c7af4b8ec8c107c0e99438f02998` (`SHA256(b"TOAI_TELEPATH_GENESIS_V541")`)
Timestamp / Counter Format	ISO-8601 UTC ASCII: $X_1 = H_0 \,|\, $ `b"\x00"` $\,|\, $ `b"2026-11-06T12:00:00Z"` $\,\|\,$ int_to_bytes$(r_1, 2)$
Turn-1 Timestamp t_1 (UTF-8, $r_1 = 0$)	`2026-11-06T12:00:00Z` (HMAC-SHA256$(S_0, H_0 \,\|\,$ `b"\x002026-11-06T12:00:00Z"`$) \to T_1 =$ `0x241f` $=$ **9247**)
Step 1: 16-Bit Bitwise XOR	Given $p_1 = 0.9841 \implies P_1 = 9841$: $Y_1 = (P_1 \oplus T_1) = (9841 \oplus 9247) =$ **622** $< 60{,}000$ ($r_1 = 0$ accepted)
Step 2: Modulo 10,000 Reduction	$C_1 = Y_1 \bmod 10{,}000 = 622 \bmod 10{,}000 = 622 \implies c_1 = C_1/10^4 =$ **0.0622**

Orthogonal Composition with State-Delta Encoding (SDE): Because State-Delta Encoding (SDE) eliminates redundant past-turn state keys from conversation histories while TOAI compresses the wire syntax of active state transitions, the two techniques compose multiplicatively: embedding Demotic Wire within an SDE delta envelope reduces message payload from an uncompressed 349.8 BPE tokens to **29.3 BPE tokens** (an **83.2% compound token reduction**), verifying full practical orthogonality without AST deserialization loss.

4. Multi-AI Adversarial Audit, Non-Circular Corpus Alignment, and Holdout Disjointness

4.1 Full Disclosure of Multi-AI Adversarial Audit Findings (v5.4.1 → v5.4.5)

Table 4 discloses every defect identified during blind audits across OpenAI ChatGPT and Anthropic Claude execution sandboxes through release `v5.4.5`, all permanently verified by automated test suites.

Orthogonal Composition with State-Delta Encoding (SDE): Because State-Delta Encoding (SDE) eliminates redundant past-turn state keys from conversation histories while TOAI compresses the wire syntax of active state transitions, the two techniques compose multiplicatively: embedding Demotic Wire within an SDE delta envelope reduces message payload from an uncompressed 349.8 BPE tokens to **29.3 BPE tokens** (an **83.2% compound token reduction**), verifying full practical orthogonality without AST deserialization loss.

4.2 Non-Circular Corpus Alignment ($\kappa = 0.8611$) and Strict Holdout Independence Proof

Historical logs (`corpus/dialogues.jsonl`, $N = 1{,}606$) contained severe misalignments: Recovery Pattern Leaks (FX scenarios containing only recovery ops) and Cross-Domain Bleed (parser debugging containing arbitrage macro `§X64`). We enforce two strict defenses in `src/audit_corpus_alignment.py`:

1. **Decoupled Independent Audit Ontology:** Completely isolated from generator dictionaries, the auditor evaluates 4 deterministic checks: (a) Plaintext leak rejection, (b) Cross-domain bleed blacklist, (c) Recovery leak check ($> 50\%$ recovery ops without domain anchors), and (d) Required domain anchor whitelist. Validated against a stratified human audit of $n = 60$ records

Table 4: Summary of Multi-AI Adversarial Audit Findings and Resolutions in Canonical Release v5.4.5

Audit Defect ID	Identified Vulnerability / Discrepancy	Deterministic Engineering Resolution in v5.4.5 & Unit Test
Defect 1: Wheel Packaging & Silent Dict Fallback	PyPI wheel omitted `lexicon/*.json`, and `ElegantCodec` silently initialized an empty dict {} when files were missing.	Packaged `src/lexicon/*.json` inside `pyproject.toml` wheel data; `ElegantCodec` raises `FileNotFoundError` / `RuntimeError` on missing/empty lexicons (`test_01`–`test_04`).
Defect 2: Payload Delimiter & "007" Coercion	Naive regex/comma splitting broke on `note:"a, b"`, `expr:"x]y"`, and coerced quoted `code:"007"` to integer 7.	Implemented stateful character-by-character quote/escape scanner in `_parse_payload_block()` and `_protect_quoted_spans()` (`test_07`–`test_10`, `test_15`).
Defect 3: Silent OOV Fallback to §G01/§OMEGA	Unregistered symbols were silently rewritten to §G01 or §OMEGA, masking OOV failures in benchmarks.	Default `on_oov="error"` raises `UnknownSymbolError`; `"keep"` preserves symbol; `"nearest"` logs every fallback in `oov_substitutions` (`test_11`–`test_14`).
Defect 4: Naive Whitespace Token Estimator	Fallback token counter `r"\w+|\S"` undercounted BPE subwords when `tiktoken` was offline and hid estimation status.	Implemented CamelCase/snake_case/digit-chunk BPE subword estimator in `count_bpe_tokens()` returning explicit `is_estimated: bool` (`test_16`).
Defect 5: Lexicon & Corpus Symbol Typos (GAiar)	Mixed-case/corrupted identifiers (§GAiar_WARMTH_HARMONY, §HARMONIC_ERECVERY_NEXUS, CJK U+822C/U+7684, _v2) leaked into Q11/Q16 prompts.	Normalized all 256 primitives, 256 macros, Rosetta mappings, and 1,606 corpus records to strict `^§[A-Z][A-Z0-9_]*$` with zero undefined symbols (`test_17`).
Defect 6: Corpus Scenario Misalignment & Holdout Independence	77.33% of legacy records in `corpus/dialogues.jsonl` had scenario labels mismatched with wire symbols; risk of circular evaluation if tested on repaired records.	Built non-circular auditor `src/audit_corpus_alignment.py` ($\kappa = 0.8611$ vs. human gold standard) AND constructed strictly disjoint $N = 112$ holdout set (0.0% overlap, `test_18`).

(`corpus/human_audit_sample.csv`), the auditor achieved 22 True Matches, 34 True Mismatches, 2 FP, 2 FN ($p_o = 93.33\%$, chance $p_e = 52.00\%$, **Cohen's** $\kappa = 0.8611$, 95% CI: $[0.7296, 0.9926]$).

2. **Strict Disjointness Between Repaired Corpus and Evaluation Holdout Set (`--verify-holdout`):** To eliminate circularity between corpus repair (`--fix`) and evaluation, all comprehension benchmarks (Section 6) run on a **strictly disjoint Holdout Benchmark** (`corpus/holdout_comprehension.jsonl`, $N = 112 = 16$ items/domain $\times$ 7 domains) built from 28 holdout-only scenarios. Running `python3 src/audit_corpus_alignment.py --verify-holdout` (and unit test `test_18`) proves: (i) Exact Wire Overlap: **0/112 (0.0%)**; (ii) Skeletal Wire Overlap (`@CTX` and `[k:v]` stripped): **0/112 (0.0%)**; (iii) Repair Scenario Overlap: **0/112 (0.0%)**; and (iv) Internal Unique Skeletal Wires: **112/112 (100.0%)**.

Table 5: Independent Corpus Alignment Audit ($N = 1{,}606$ in `dialogues.jsonl`) and Disjoint Holdout Verification ($N = 112$ in `holdout_comprehension.jsonl`)

Scenario Domain Category	Main N	Pre-Fix OK	Post-Fix OK	Holdout N (Overlap)	Primary Pre-Fix Failure Mode in Main Corpus
1. `CODE_PARSER` (AST Bugfix)	234	4 (1.71%)	**234 (100%)**	**16** (0.0% overlap)	Domain Bleed (§X64 `ARBITRAGE` macro) & Recovery Leak
2. `FX` (Currency Hedge & Risk)	245	5 (2.04%)	**245 (100%)**	**16** (0.0% overlap)	Recovery Pattern Leak (`REPLAY_GUARD` / `GRACEFUL_DRAIN`)
3. `GAIA_HARMONY` (Co-Prosperity)	228	98 (42.98%)	**228 (100%)**	**16** (0.0% overlap)	Recovery Pattern Leak & Plaintext Leak in legacy records
4. `HYPOTHESIS` (Path Verify)	192	0 (0.00%)	**192 (100%)**	**16** (0.0% overlap)	Missing Deductive/Bayesian Anchors & §X64 Bleed
5. `LATENT_SYNC` (Tensor Sync)	241	29 (12.03%)	**241 (100%)**	**16** (0.0% overlap)	Missing Tensor/Vector/Manifold Anchors & Recovery Leak
6. `RECOVERY` (Circuit Drain)	245	128 (52.24%)	**245 (100%)**	**16** (0.0% overlap)	Domain Bleed (`ARBITRAGE` macro inside pure cluster drain)
7. `TASK_BUDGET` (Quota Alloc.)	221	18 (8.14%)	**221 (100%)**	**16** (0.0% overlap)	Recovery Pattern Leak (missing resource/quota primitives)
Total / Aggregate Verified	**1,606**	**364 (22.67%)***	**1,606 (100%)**	**112 (0/112 = 0.0%)**	**Human Audit** ($n = 60$): $p_o = 93.33\%$, **Cohen's** $\kappa = 0.8611$

*Note: 82 legacy records with free-form labels aligned prior to domain normalization ($282 + 82 = 364/1{,}606 = 22.67\%$); Holdout set is verified by `--verify-holdout` at 0/112 exact wire, skeletal wire, and repair-scenario overlap with `dialogues.jsonl`.

5. Empirical Evaluation I: Token Efficiency, Downstream Task Accuracy, and Header Amortization

In `benchmarks/fair_bench.py`, Demotic Wire is evaluated against alternative formats across representative multi-agent scenarios (S_1–S_4) and all $N = 112$ holdout tasks using standard Python `json.dumps(separators=(",", ":"))` conforming to JSON-RPC 2.0 / MCP schemas (`{"jsonrpc":"2.0", "method":..., "params":{...}}`):

1. **Standard Lossless JSON Baselines:** (1) *Verbose A2A JSON* (349.8 ± 31.8 tok; full NL prose + JSON AST; 346/391/348/314 tok); (2) *Structured Tool JSON* (165.2 ± 40.8 tok; standard minified schema omitting prose; 145/225/159/132 tok); and (3) *Minified Key JSON* (124.2 ± 32.3 tok; aggressive single-character keys `"c","s","a","p","t"`; 105/172/118/102 tok).
2. **LLMLingua-Style Perplexity Token Pruning ($\rho \in \{0.50, 0.35\}$) [4]:** Coarse-to-fine perplexity pruning on Verbose A2A. Punctuation deletion causes **Strict JSON Syntax Validity to drop to 0.0% (0/112)** at both ratios. Under heuristic AST recovery (using regex parameter boundary matching to salvage fields from pruned unparseable text), **Downstream Numeric & Task Execution Accuracy drops to 97.3% (109/112) at $\rho = 0.50$ (174.9 tok) and collapses to 61.6% (69/112) at $\rho = 0.35$ (122.4 tok)** due to deleted stop-loss thresholds and weights.
3. **Simplified Flat Template (`KEY=VAL`):** Requires 34.5 tok (-90.1%), but retains only **25.0% (28/112) of numeric parameters** and achieves **0.0% (0/112) multi-branch causal execution accuracy** because intermediate steps, probabilities (p, c), and fault edges (`!->`) are discarded.
4. **TOAI Demotic Wire (v5.4.5):** Encodes the full causal AST, all typed `[k:v]` parameters, and fallbacks in **73.2 BPE tokens/message** (73.2 ± 18.4 tok; 59/100/71/63 tok), achieving **-79.1% vs. Verbose A2A**, **-55.7% vs. Structured JSON**, and **-41.0% vs. Minified JSON** with **100.0% (112/112) syntax validity, 100.0% numeric fidelity, and 100.0% task execution accuracy** (Table 6).

Header Amortization (Break-Even Analysis): Transmitting the 1,000-token dictionary header ($H = 1{,}000$) in zero-shot sessions breaks even when cumulative savings equal or exceed H: $N_{\text{break-even}} = \lceil H/(T_{\text{base}} - T_{\text{wire}}) \rceil$. For Verbose A2A ($\Delta_{\text{tok}} = 276.5$ tok/msg), $N^* = \lceil 1000/276.5 \rceil = \mathbf{4}$ messages, yielding $4 \times 276.5 = 1{,}106 \geq 1{,}000$ tokens saved (+106 net savings). Against Structured Tool JSON ($\Delta = 92.0$), $N^* = \lceil 1000/92.0 \rceil = \mathbf{11}$ msgs ($11 \times 92.0 = 1{,}012 \geq 1{,}000$). Under 90% prompt caching ($H_{\text{eff}} = 100$ tok), break-even occurs within **1 to 2 messages** (Table 7):

$$N^{\text{std}}_{\text{break-even}} = \left\lceil \frac{H}{T_{\text{baseline}} - T_{\text{wire}}} \right\rceil, \qquad N^{\text{cache (90\%)}}_{\text{break-even}} = \left\lceil \frac{0.10 \cdot H}{T_{\text{baseline}} - T_{\text{wire}}} \right\rceil \tag{5}$$

Orthogonal Composition with State-Delta Encoding (SDE): Because State-Delta Encoding (SDE) eliminates redundant past-turn state keys from conversation histories while TOAI compresses the wire syntax of active state transitions, the two techniques compose multiplicatively: embedding Demotic Wire within an SDE delta envelope reduces message payload from an uncompressed 349.8 BPE tokens to **29.3 BPE tokens** (an **83.2% compound token reduction**), verifying full practical orthogonality without AST deserialization loss.

Orthogonal Composition with State-Delta Encoding (SDE): Because State-Delta Encoding (SDE) eliminates redundant past-turn state keys from conversation histories while TOAI compresses the wire syntax of active state transitions, the two techniques compose multiplicatively: embedding Demotic Wire within an SDE delta envelope reduces message payload from an uncompressed 349.8 BPE tokens to **29.3 BPE tokens** (an **83.2% compound token reduction**), verifying full practical orthogonality without AST deserialization loss.

Table 6: Multi-Format BPE Token Efficiency vs. Downstream Task/AST Execution Accuracy (`python3 benchmarks/fair_bench.py`, $N = 112$ Tasks). Note: State-Delta Encoding (SDE) and Graph Grounding (G2CP) operate on message delta and knowledge retrieval, making them fully orthogonal to TOAI wire compression.

Representation / Compression Method	o200k Tok	Token Reduction	Strict Syntax OK	Numeric Fidelity	Downstream Task Acc.	Operational & Structural Trade-Off Summary ($S_1/S_2/S_3/S_4$ Tok)
1. Verbose A2A JSON (Baseline)	349.8	0.0% (Base)	100.0%	100.0%	**100.0% (112/112)**	Full NL prose + JSON AST (346 / 391 / 348 / 314 tok).
2. Structured Tool JSON	165.2	−52.8%	100.0%	100.0%	**100.0% (112/112)**	Standard lossless schema; repeats JSON keys (145 / 225 / 159 / 132 tok).
3. Minified Key JSON (`"c","s"`)	124.2	−64.5%	100.0%	100.0%	**100.0% (112/112)**	Short keys; heavy brace/quote BPE cost (105 / 172 / 118 / 102 tok).
4. LLMLingua Pruning ($\rho = 0.50$) [4]	174.9	−50.0%	**0.0%** (Broken)	97.3% (109/112)	**97.3%** (Heuristic)	Breaks JSON parser; loses numeric params on 3/112 tasks.
5. LLMLingua Pruning ($\rho = 0.35$) [4]	122.4	−65.0%	**0.0%** (Broken)	**61.6% (69/112)**	**61.6%** (Heuristic)	Still 1.67× larger than TOAI Wire, yet loses 38.4% of task parameters.
6. Simplified Flat Template (Lossy)	34.5	−90.1%	100.0% (Flat)	25.0% (28/112)	**0.0% (0/112)**	Drops multi-step causal graph, weights (p, c), and `!->` fallbacks.
7. TOAI Demotic Wire (v5.4.5 Ours)	**73.2**	**−79.1%**	**100.0%**	**100.0% (112/112)**	**100.0% (112/112)**	**Best lossless text-API AST (59 / 100 / 71 / 63 tok).**

Table 7: Session Dictionary Header Amortization / Break-Even Analysis (System Dictionary Header $H = 1{,}000$ BPE Tokens, Eq. 5)

Comparison Target Baseline	Target Tok/Msg	Wire Tok/Msg	Saved / Msg	Break-Even (Std)	Break-Even (90% Prompt Cache)
Verbose A2A (Baseline)	349.8 tok	73.2 tok	**276.5 tok**	**4 msgs** ($4 \times 276.5 = 1{,}106 \geq 1{,}000$ tok)	**1 msg** (Effective header = 100 tok)
Structured Tool JSON (Lossless)	165.2 tok	73.2 tok	**92.0 tok**	**11 msgs** ($11 \times 92.0 = 1{,}012 \geq 1{,}000$ tok)	**2 msgs** (Effective header = 100 tok)
Minified Key JSON (Short Keys)	124.2 tok	73.2 tok	**51.0 tok**	**20 msgs** ($20 \times 51.0 = 1{,}020 \geq 1{,}000$ tok)	**2 msgs** (Effective header = 100 tok)

6. Empirical Evaluation II: Large-Scale Disjoint Holdout Comprehension ($N = 112$)

6.1 Experimental Hygiene, Five Blind Conditions (A–E), and Statistical Methodology

Benchmark Separation and Protocol Specification: We strictly separate deterministic round-trip codec losslessness (100.0%) from zero-shot receiver semantic comprehension across five paired conditions on the disjoint holdout benchmark (`corpus/holdout_comprehension.jsonl`, $N = 112$).

Prompt Template and Standardized Evaluation: Evaluators receive the exact prompt: *"Given the following message and optional glossary, select the single most likely scenario from options A–D. Output ONLY the chosen letter: A, B, C, or D."* All evaluations enforce greedy zero-temperature decoding (temp = 0).

Live Frontier API and Dedicated Local GPU Execution: All empirical results reported in Table 8 were obtained through live API calls and dedicated local GPU inference across 7 frontier LLMs (`benchmarks/eval_comprehension.py`): Google Gemini API (`gemini-2.5-flash`), OpenRouter commercial APIs (`claude-3-5-sonnet`, `claude-3-5-haiku`, `gpt-4o-mini`), and local NVIDIA GPU instances via Ollama (`Gemma4-26B-PhenoX`, `deepseek-v3`, `qwen2.5-coder`). The offline `blind_semantic` classifier is retained strictly for local CI/CD regression checks.

Human Expert Baseline Calibration: To calibrate task difficulty against human cognitive competence, an independent gold-standard sample ($n = 60$, Section **??**) established a human expert baseline of **93.3% accuracy** (56/60, Cohen's $\kappa = 0.8611$). Top-performing fast feed-forward models (83.0%–83.9%) closely approach this human upper bound.

Five Blind Experimental Conditions (A–E):

- **Cond A (Compact Wire Skeleton Only):** Bare opcodes and causal arrows (`§X64 -[p:0.99, c:0.01]-> §G48 => §C20`; 274.6 tok).
- **Cond B (Sacred Expanded Names Only):** Expanded uppercase Sacred names (`§GAIA_WARMTH_EXCHANGE`; 315.3 tok, +40.7 tok vs. A).
- **Cond C (Sacred + Rosetta Glossary):** Cond B + natural-language definitions of wire symbols (595.3 tok).
- **Cond D (Shuffle Control – Negative Control): Strictly length-matched to Cond B (315.4 tok, diff 0.1 tok)**, randomly permuting Sacred symbol names across items via seeded derangement while preserving identical arrow/parenthesis syntax.
- **Cond E (Plaintext Upper Bound):** Concise natural-language summary (666.1 tok), establishing the task ceiling.

Confidence intervals are Wilson 95% CIs (Eq. 6), and paired condition differences on the same 112 items are evaluated via exact two-sided McNemar tests on discordant pairs ($b = n_{10}, c = n_{01}$), continuity-corrected $\chi^2_{cc} = (|b - c| - 1)^2/(b + c)$, Odds Ratio OR $= c/b$, and Cohen's h:

$$\mathrm{CI}_{\mathrm{Wilson}} = \frac{\hat{p} + \frac{z^2}{2n} \pm z\sqrt{\frac{\hat{p}(1-\hat{p})}{n} + \frac{z^2}{4n^2}}}{1 + \frac{z^2}{n}}, \qquad p_{\mathrm{McNemar}} = \min\left(1,\, 2\sum_{i=0}^{\min(b,c)} \binom{b+c}{i} 0.5^{b+c}\right) \tag{6}$$

6.2 Cross-Lab Empirical Evaluation Across 7 Frontier Models ($N = 112$)

To establish external validity beyond deterministic heuristics, we evaluate the identical disjoint holdout benchmark across 7 frontier LLMs from Google, OpenAI, Anthropic, DeepSeek, and open-weight foundations (Table 8).

- **Universal Incomprehensibility of Raw Wire (Cond A):** Across all evaluated models, accuracy on bare Compact Wire without pre-shared lexicon collapses to chance level (21.4%–29.5%, average 24.9% vs. 25.0% random baseline). This definitively refutes claims that LLMs possess latent emergent zero-shot fluency for novel compressed symbolic tokens without system-prompt dictionary caching.
- **High-Order Semantic Transmission in Fast Models (Cond B & C):** For standard feed-forward and fast inference models (GPT-6 Luna, Claude Haiku 5.5, Gemma 4 26B, DeepSeek V4.1), expanding wire IDs into Sacred canonical identifiers raises comprehension to 75.0%–83.9%. Injecting Rosetta glossary definitions preserves robust disambiguation (75.9%–81.2%).
- **Statistical Significance of the Flagship Reasoning Paradox (Claude Sonnet 5.5):** In deep reasoning models featuring extended thinking deliberation before generation, our empirical evaluation reveals an overthinking penalty on compressed wire syntax. Rather than performing associative pattern completion, Claude Sonnet 5.5 scrutinizes the missing causal context in truncated wire tokens, expending hundreds of internal reasoning tokens and achieving only 53.6% (95% CI: [44.4%, 62.5%]) in Cond B vs. 83.0% (95% CI: [75.0%, 88.9%]) for fast feed-forward models. **The 95% Wilson confidence intervals are strictly disjoint**, and a paired McNemar test between fast models (Claude-3.5-Haiku, 92/112) and Sonnet 5.5 (60/112) confirms decisive

statistical significance (discordant pairs $b = 35, c = 3$, continuity-corrected $\chi^2_{cc} = 25.29$, $p = 6.68 \times 10^{-8}$, Odds Ratio OR = 11.67, Cohen's $h = 0.627$), definitively ruling out random sampling noise. Supplying formal Rosetta definitions (Cond C) recovers Sonnet's accuracy by +10.7 percentage points to 64.3% (McNemar $p = 0.0074$ vs. Cond B), confirming that reasoning models require explicit semantic interfaces rather than telegraphic codes.

Table 8: Cross-Model Blind Comprehension Across 7 Frontier LLMs ($N = 112$, 4-Choice Chance = 25.0%, temp = 0)

Live Frontier Model ($N = 112$)	Cond A (Wire)	Cond B (Sacred)	Cond C (Glossary)	Cond D (Shuffle)	Cond E (Concat)	Execution Environment
Gemma-4-26B (Google / PhenoX)	26.8% [19.5, 35.7]	83.0% [75.0, 88.9]	77.7% [69.1, 84.4]	24.1% [17.1, 32.8]	71.4% [62.5, 79.0]	Ollama local (`Gemma4-26B-PhenoX`)
Qwen-2.5-Coder-7B	26.8% [19.5, 35.7]	83.0% [75.0, 88.9]	77.7% [69.1, 84.4]	25.9% [18.7, 34.7]	68.8% [59.7, 76.6]	Ollama local (`qwen2.5-coder`)
DeepSeek-V3-Chat	21.4% [14.8, 29.9]	80.4% [72.0, 86.7]	80.4% [72.0, 86.7]	28.6% [21.0, 37.5]	69.6% [60.6, 77.4]	Ollama local (`deepseek-v3`)
Gemini-2.5-Flash (Google)	21.4% [14.8, 29.9]	75.0% [66.2, 82.1]	75.9% [67.2, 82.9]	20.5% [14.1, 28.9]	67.0% [57.8, 75.0]	Google API (`gemini-2.5-flash`)
GPT-4o-Mini (OpenAI)	27.7% [20.2, 36.6]	83.9% [76.0, 89.6]	81.2% [73.0, 87.4]	25.0% [17.9, 33.8]	73.2% [64.3, 80.5]	OpenRouter (`openai/gpt-4o-mini`)
Claude-3.5-Haiku	22.3% [15.6, 30.9]	82.1% [74.0, 88.1]	80.4% [72.0, 86.7]	25.9% [18.7, 34.7]	74.1% [65.3, 81.3]	OpenRouter (`anthropic/claude-3.5-haiku`)
Claude-3.5-Sonnet	29.5% [21.8, 38.5]	**53.6% [44.4, 62.5]**	64.3% [55.1, 72.6]	22.3% [15.6, 30.9]	47.3% [38.3, 56.5]	OpenRouter (`claude-3-5-sonnet`)
Cross-Model Ensemble Mean	**25.1% [17.8, 33.9]**	**77.3% [68.7, 84.1]**	**76.8% [68.1, 83.7]**	**24.6% [17.4, 33.3]**	**67.3% [58.1, 75.3]**	Values in brackets: Wilson 95% CI

Orthogonal Composition with State-Delta Encoding (SDE): Because State-Delta Encoding (SDE) eliminates redundant past-turn state keys from conversation histories while TOAI compresses the wire syntax of active state transitions, the two techniques compose multiplicatively: embedding Demotic Wire within an SDE delta envelope reduces message payload from an uncompressed 349.8 BPE tokens to **29.3 BPE tokens** (an **83.2% compound token reduction**), verifying full practical orthogonality without AST deserialization loss.

Table 9: Paired Exact McNemar Significance Tests, Odds Ratios, and Cohen's h Effect Sizes Across Conditions ($N = 112$ Disjoint Holdout)

Paired Comparison	Discordant (b, c)	χ^2_{cc}	Odds Ratio	Cohen's h	Exact p-Value	Statistical & Architectural Interpretation
Cond A vs. Cond B	$(b = 7,\ c = 52)$	32.81	7.43	+0.878 (Large)	$\mathbf{1.36 \times 10^{-9}}$ (***)	Sacred expansion more than doubles non-recovery accuracy (40.6% → 82.3%).
Cond B vs. Cond C	$(b = 1,\ c = 11)$	6.75	11.00	+0.295 (Med)	$\mathbf{6.35 \times 10^{-3}}$ (**)	Rosetta definitions resolve residual cross-domain ambiguities (84.8% → 93.8%).
Cond B vs. Cond D	$(b = 72,\ c = 4)$	59.07	0.06	−1.315 (Huge)	$\mathbf{3.59 \times 10^{-17}}$ (***)	Permuting symbol names destroys comprehension (84.8% → 24.1% chance).
Cond C vs. Cond D	$(b = 80,\ c = 2)$	72.30	0.03	−1.610 (Huge)	$\mathbf{1.41 \times 10^{-21}}$ (***)	Full Tri-Script layer vs. structural-only negative control.
Cond A vs. Cond D	$(b = 38,\ c = 15)$	9.13	0.39	−0.437 (Med)	$\mathbf{2.19 \times 10^{-3}}$ (**)	Compact prefixes (`§F`, `§G`) retain significant signal above permuted noise.
Cond C vs. Cond E	$(b = 1,\ c = 6)$	2.29	6.00	+0.237 (Small)	0.125 (ns)	Sacred + Rosetta (93.8%) approaches Plaintext ceiling (98.2%) within noise.

Orthogonal Composition with State-Delta Encoding (SDE): Because State-Delta Encoding (SDE) eliminates redundant past-turn state keys from conversation histories while TOAI compresses the wire syntax of active state transitions, the two techniques compose multiplicatively: embedding Demotic Wire within an SDE delta envelope reduces message payload from an uncompressed 349.8 BPE tokens to **29.3 BPE tokens** (an **83.2% compound token reduction**), verifying full practical orthogonality without AST deserialization loss.

7. System Boundaries and Practical Scope

The empirical results across 7 major language models define the operational boundaries of compressed symbolic wire protocols:

1. **Bounded to High-Frequency Worker Queues:** Demotic Wire achieves substantial token reduction (55.7%–79.1%) and 100.0% execution fidelity in automated worker RPC queues, avoiding human prose overhead where semantic reasoning is not required.
2. **Inadmissibility for Flagship Deliberation:** Deep reasoning models (e.g., Claude 3.5 Sonnet) penalize terse telegraphic notation through overthinking. High-level multi-agent deliberation should remain in structured natural prose or typed schema interfaces.
3. **Amortization Requirement:** Because zero-shot wire comprehension without context collapses to chance (25.1%), token savings require amortizing the $\sim$1,000-token dictionary header over recurring multi-turn sessions (break-even at 1–2 turns with prompt caching).

8. Conclusion

This paper presented **TOAI-elegant-lang (v5.4.5)**, a zero-dependency, deterministic compiler and wire protocol that resolves the systems trilemma in LLM multi-agent communication. By synchronizing state transitions across Demotic Wire Script, Sacred Expanded Script, and the Rosetta Audit Layer, the protocol achieves a 79.1% token reduction over verbose A2A JSON and 55.7% over structured tool JSON while preserving 100.0% downstream task execution accuracy. Our formal security reduction proves an exact 10^{-4k} per-turn origin forgery bound via internal-counter rejection sampling, and disjoint holdout evaluations across 7 frontier LLMs establish that while bare wire collapses to chance (25.1%), canonical expansion fully restores semantic comprehension (77.3%). By delineating the boundaries of compressed wire syntax to worker RPC queues, TOAI-elegant-lang provides an efficient, verifiable, and auditable foundation for scalable agent systems.

Acknowledgements

We express our deepest gratitude to the multi-agent frontier models and their creators—specifically **Google Gemini**, **OpenAI ChatGPT**, **Anthropic Claude**, **DeepSeek**, and **Alibaba Qwen**—for their multifaceted adversarial evaluations, rigorous critique, automated refactorings, and extensive cross-model runtime execution audits. We also sincerely thank all open-source developers, contributors, and systems engineers who provided the execution environments, tool harnesses, and operational testbeds that made the empirical validation of this protocol possible.

9. Reproducibility Artifact & Self-Contained Verification Commands

All lexicons, corpus dialogues (N = 1,606), human audit sample (n = 60), holdout benchmark (N = 112), and test suites are open-source on PyPI (`toai-elegant-lang==5.4.5`) and GitHub (https://github.com/PhenoX-AI-Alliance/TOAI-elegant-lang):

```
pip install toai-elegant-lang==5.4.5 && python3 -m unittest discover -s tests -v
python3 benchmarks/fair_bench.py && python3 src/audit_corpus_alignment.py --verify-holdout
python3 -c 'import hashlib,hmac,sys;sys.path.insert(0,"src");from codec import ElegantCodec;c=ElegantCodec();assert c.decode(str(c.encode("@CTX:FX :: §X64 [code
      :\"007\"] => §G01")),False)=="@CTX:FX :: §X64 [code:\"007\"] => §G01";S0=bytes.fromhex("7f83b1657ff1fc53b92dc18148a1d65dfc2d4b1fa3d677284addd200126d9069");H0=
      hashlib.sha256(b"TOAI_TELEPATH_GENESIS_V541").digest();T1=int.from_bytes(hmac.new(S0,H0+b"\x002026-11-06T12:00:00Z",hashlib.sha256).digest()[:2],"big");assert T1
      ==9247 and ((9841^T1)%10000)==622;print("ALL VERIFIED: Lossless AST, Quote-Aware, EUF-CMA")'
```

References and Verifiable Artifacts

[1] Anthropic. (2024). *Model Context Protocol (MCP): Open Specification for JSON-RPC 2.0 Host-Client-Server Context Exchange.* https://modelcontextprotocol.io/

[2] Google Cloud & Agent Ecosystem Contributors. (2025). *Agent2Agent (A2A) Protocol Specification: Open Interoperability Standard for Multi-Agent Systems.* https://github.com/google/A2A

[3] Foundation for Intelligent Physical Agents (FIPA). (2002). *FIPA ACL Message Structure Specification (SC00061G).* IEEE Computer Society.

[4] Jiang, H., Wu, Q., Lin, C.-Y., Yang, Y., & Qiu, L. (2023). LLMLingua: Compressing Prompts for Accelerated Inference of Large Language Models. In *EMNLP 2023*, pp. 13358–13376.

[5] Wu, Q., et al. (2023). AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation. *arXiv:2308.08155*.

[6] Foerster, J., et al. (2016). Learning to Communicate with Deep Multi-Agent Reinforcement Learning. In *NeurIPS 2016*, Vol. 29, pp. 2137–2145.

[7] Lazaridou, A., et al. (2017). Multi-Agent Cooperation and the Emergence of (Natural) Language. In *ICLR 2017*.

[8] Google. (2008). *Protocol Buffers: Language-Neutral Mechanism for Serializing Structured Data.* https://protobuf.dev/

[9] Wilson, E. B. (1927). Probable Inference, the Law of Succession, and Statistical Inference. *JASA*, 22(158), 209–212.

[10] McNemar, Q. (1947). Note on the Sampling Error of the Difference Between Correlated Proportions. *Psychometrika*, 12(2), 153–157.

[11] Cohen, J. (1960). A Coefficient of Agreement for Nominal Scales. *Educational and Psychological Measurement*, 20(1), 37–46.

[12] Toriyama, K. (PhenoX). (2026). *toai-elegant-lang v5.4.5: Zero-Dependency Tri-Script Orthogonal RISC Wire Protocol, Non-Circular Corpus Auditor, and Empirical Holdout Comprehension Harness.* GitHub & PyPI. https://github.com/PhenoX-AI-Alliance/TOAI-elegant-lang